

worst, best and average case are all recorded and reported when a search is done, or when all dam info is printed. The *Dam* class implements the Comparable interface (used by both the BST and AVL), and stores values for name, dam level and FSC. I used the *toString()* method inside the *Dam* class for easy printing of information during search. The *printDamName()* function is implemented in both the *damAVLApp* and *damBSTApp*. Both the *damBSTApp* and *damAVLApp* search through the nodes (containing *Dam* objects) and compare the query dam name - for which a temporary *Dam* object is created - to the names of each node and moves down the branches according to lexicographic ordering (*string.compareTo(otherString)*) until the queried Dam is found, or the end of a branch is reached. In both applications the *opCount* starts at zero when each search function is initiated, and is incremented by 1 each time a comparison is made - I used an int array to take advantage of pass-by-reference (as with the *opCounts* for insertion). I have written bash scripts to run both search and *printAllDam()* functions with test values, redirecting the output to text files (using *echo (function()) > testFile.txt*).

For the main experiment - part 5 - plotting the comparison between *opCounts* when inserting new nodes and searching varied subsets of data - my method was as follows:

- Create subsets of the original data file (from a set of 1, to 211), using a bash script (a loop and the unix head command, writing results to 211 new CSV files, *test_1.txt... test_211.txt*)
- Create *test()* functions in both the AVL and binary search tree applications to loop through each subset of the data, create a temporary *damAVLApp* or *damBSTApp* and create a *dataFrame* in the object containing the subset of data.
- The *opCount* statistics for insertion are recorded for each data subset and added to a string-builder, in order to be written to file.
- Then search once for each name in the given subset, storing the operation count for each search in an array.
- For a given subset, the array of operation counts will have a max and min value, which are recorded as the worst and best cases for that dataset. The average is also calculated.
- Finally, resulting curves are plotted (best, average and worst) on a two graphs, one for the AVL search algorithm, one for the binary search tree search algorithm for both search and insertion *opCount* statistics using the *graph* class (see part d.)

c. Test Results: Part 2 and 4

part 2: results from the Binary Search Tree Application

Test 1 (part 1 and 2)

Searching BST for: 3 known dam names

Result:

name: Lindleyspoort Dam

dam level: 2.7

FSC: 14.208

the size of the dataset was: 211

The number of insertion comparisons (best, average and worst) were: 0, 7, 12

The total number of search comparisons was: 3

Result:

name: Koppies Dam

dam level: 39

FSC: 42.311

the size of the dataset was: 211

The number of insertion comparisons (best, average and worst) were: 0, 7, 12

The total number of search comparisons was: 6

Result:

name: Woodstock Dam

dam level: 77.099999999999994

FSC: 373.25

the size of the dataset was: 211

The number of insertion comparisons (best, average and worst) were: 0, 7, 12

The total number of search comparisons was: 9

Test 2 (part 1 and 2)

Searching BST for: 1 unknown dam name

Result:

Dam not found

query: Nothing

the size of the dataset was: 211

The number of insertion comparisons (best, average and worst) were: 0, 7, 12

The total number of search comparisons was: 8

Test 3 (part 4)

First 10 entries:

Printing all dam entries from BST
(Pre_order Traversal)

Result:

name: Ngotwane Dam
level: 4.6
FSC: 19.0330000000000001

name: Hartbeespoort Dam
level: 96.5
FSC: 186.44

name: Bon Accord Dam
level: 103
FSC: 4.381

name: Albasini Dam
level: 69.2
FSC: 28.1990000000000002

name: Blyderivierpoort Dam
level: 50
FSC: 54.369

Test 3 (part 4)

name: Allemanskraal Dam
level: 13.5
FSC: 174.52

name: Albert Falls Dam
level: 36.4
FSC: 288.14

name: Bloemhof Dam
level: 15.5
FSC: 1240.24

name: Armenia Dam
level: 10
FSC: 12.9570000000000001

name: Berg River Dam
level: 31.1
FSC: 127.05

...

Test 3 (part 4) (continued)

Last 10 entries:

...

name: Waterdown Dam
level: 90.1
FSC: 37.4410000000000003

name: Wemmershoek Dam
level: 52.3
FSC: 58.71

name: Westoe Dam
level: 48.3
FSC: 60.094999999999999

name: Wolwedans Dam
level: 95.4
FSC: 24.6260000000000001

name: Witklip Dam
level: 39.299999999999997
FSC: 12.519

...

Test 3 (part 4)

name: Xonxa Dam
level: 100
FSC: 115.86

name: Wriggleswade Dam
level: 98.7
FSC: 91.4710000000000004

name: Woodstock Dam
level: 77.099999999999994
FSC: 373.25

name: Xilinxha Dam
level: 27.2
FSC: 13.823

name: Zaaiohoek Dam
level: 58.5
FSC: 184.63

Printed all dam entries from binary
search tree.

the size of the dataset was: **211**
The number of insertion comparisons
(best, average and worst) were: **0, 7, 12**

Test 1 (part 3 and 4)

Searching AVL Tree for: 3 known dam names

Result:

name: Lindleyspoort Dam
dam level: 2.7
FSC: 14.208
the size of the dataset was: 211
The number of insertion comparisons (best, average and worst) were: 0, 7, 10
The total number of search comparisons was: 1

Result:

name: Koppies Dam
dam level: 39
FSC: 42.311
the size of the dataset was: 211
The number of insertion comparisons (best, average and worst) were: 0, 7, 10
The total number of search comparisons was: 6

Result:

name: Woodstock Dam
dam level: 77.099999999999994
FSC: 373.25
the size of the dataset was: 211
The number of insertion comparisons (best, average and worst) were: 0, 7, 10
The total number of search comparisons was: 7

Test 2 (part 3 and 4)

Searching AVL Tree for: 1 unknown dam name

Result:

Dam not found
query: Nothing
the size of the dataset was: 211
The number of insertion comparisons (best, average and worst) were: 0, 7, 10
The total number of search comparisons was: 9

d. Test Results: Part 5**Test 3(part 3 and 4)**

First 10 entries:

Printing all dam entries from AVL Tree

name: Lindleyspoort Dam
level: 2.7
FSC: 14.208

name: Flag Boshielo Dam
level: 40.200000000000003
FSC: 185.13

name: Buffelspoort Dam
level: 71.400000000000006
FSC: 10.183

name: Blyderivierpoort Dam
level: 50
FSC: 54.369

name: Armenia Dam
level: 10
FSC: 12.957000000000001

Test 3(part 3 and 4)

name: Albert Falls Dam
level: 36.4
FSC: 288.14

name: Albasini Dam
level: 69.2
FSC: 28.199000000000002

name: Allemanskraal Dam
level: 13.5
FSC: 174.52

name: Berg River Dam
level: 31.1
FSC: 127.05

name: Belfort Dam
level: 98
FSC: 0.413

...

Test 3 (part 4) (continued)

Last 10 entries:

...

name: Welbedacht Dam
level: 97.5
FSC: 9.592

name: Waterdown Dam
level: 90.1
FSC: 37.441000000000003

name: Witbank Dam
level: 52.2
FSC: 104.02

name: Westoe Dam
level: 48.3
FSC: 60.094999999999999

name: Witklip Dam
level: 39.299999999999997
FSC: 12.519

Test 3 (part 4)

name: Wriggleswade Dam
level: 98.7
FSC: 91.471000000000004

name: Woodstock Dam
level: 77.099999999999994
FSC: 373.25

name: Xonxa Dam
level: 100
FSC: 115.86


name: Xilinxha Dam
level: 27.2
FSC: 13.823

name: Zaaihoek Dam
level: 58.5
FSC: 184.63

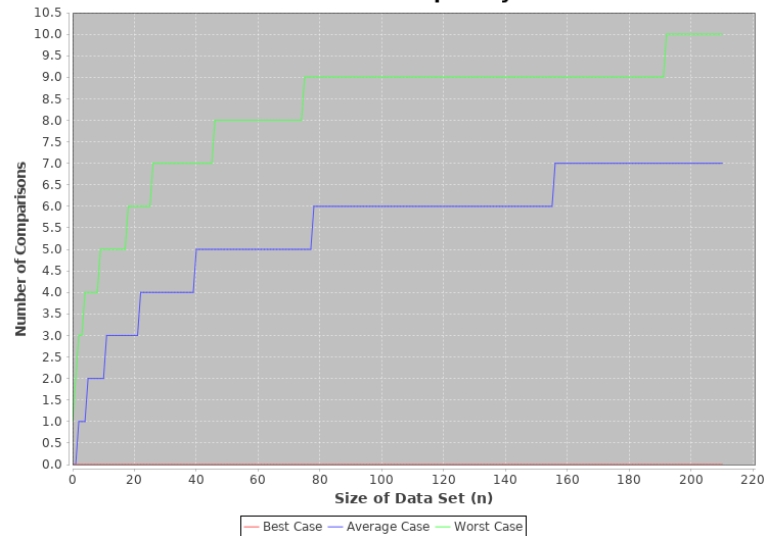
Printed all dam entries from AVL tree.
the size of the dataset was: **211**
The number of insertion comparisons
(best, average and worst) were: **0, 7, 10**

Time complexity of insert algorithms:

Binary Search Tree Insert Complexity Test




AVL Tree Insert Complexity Test




The graphs above are a graphical representation of the opCounts recorded during the insertion of data into the BS and AVL Trees, vs the size of the data set. The best, average and worst cases plotted in each case. All curves show a $\log(n)$ shape - suggesting asymptotic behaviour for large n (size dataset). By comparison the AVL tree proves to have a more efficient underlying algorithm due to the balancing of the tree during the insertion process (from the data we see an 'improvement' of anywhere between 10% and 20%, which will become increasingly significant for very large values of n). It should be noted that the AVL tree requires additional computation which we have not taken into consideration - calculating the balance factor and height of the subtrees - but this process involves integer comparisons which are significantly faster than String comparisons, which we have measured here, so we count them as negligible. It should also be noted that the best case, in every case, is 0 which is the insertion of the first piece of data at the root.

Time complexity of search algorithms:

Binary Search Tree Search Complexity Test



AVL Tree Search Complexity Test



The results from the test of time complexity of the search algorithms show similar results to that of the insertion test. The shape is, again, $\log(n)$ for all curves (roughly). By comparison, the AVL tree both increases slower (shallower gradient) in all cases, and approaches asymptotic behaviour faster than the Binary Search Tree. The advantage for the AVL Tree clearly being the fact that the data structure is balanced according to the properties of an AVL Tree: a maximum difference of 1 in the heights of subtrees. This suggests an improvement in efficiency by nearly 25% for the AVL tree, which will become very significant for large n . The correlated results of the time complexity of insert and search for both structures are correlated which is no coincidence - the binary tree shows an increase in overall complexity in the worst / average case, while the AVL Tree shows a decrease overall, which suggests balancing effects insertion time-complexity, and even more so, search time-complexity.

e. Part 6 Results

AVL search opCount Results:

best: 1
average: 6
worst: 8

AVL insert opCount Results:

best: 0
average: 7
worst: 9

BST search opCount Results:

best: 1
average: 106
worst: 211

BST insert opCount Results:

best: 0
average: 104
worst: 210

The results above are of the test run on the sorted dataset (names of Dams in alphabetical order). They are most significant in terms of displaying the power of the balancing algorithm which forms part of the insert algorithm of the AVL tree. The Binary Tree will always exhibit linear search behaviour if the data is ordered - each node connected to only one node, all on the left or right - this $O(n)$ performance is typical of an array, and very taxing computationally. The AVL Tree behaves nearly identically to when the data was unsorted.

f. Creative Components

Below I have listed what I feel are creative efforts on my part to go above and beyond the brief for the assignment.

- fixCSVFile utility class: Instead of simply using the unix head command I created a class to repair line breaks in CSV files, as well as replace commas.

- `dataFrame` class: the core of my solution to the assignment was creating a robust class to store and access information in a CSV file by column and dam name, as well as query details about its dimensions. This was used as an object in both my `DamArrayApp` and `DamBSTApp`
- I created bash scripts to run / generate test output for parts 2, 4, 5 and 6 (in the bin folder)
- `test()` methods inside `DamAVLApp` and `DamBSTApp`: I used an interesting technique to isolate each search case for single subsets of the data. By creating a temporary instance of the same object in which the `test()` method already exists (and instance within an instance of an object), I was able to effectively and cleanly isolate the set of searches run on each subset of data without creating an additional class to handle the task.

g. Git log (from most recent)

commit 979a7b76ce04c8cfe4d87befaf298100c5fc62ca
 Author: jeremy <dpljer001@myuct.ac.za>
 Date: Tue Mar 27 17:17:46 2018 +0200

javadoc 2

commit 4deaa8dfbed2333a3360729ed90659955a26b172
 Author: jeremy <dpljer001@myuct.ac.za>
 Date: Tue Mar 27 16:57:28 2018 +0200

Javadoc 1

commit 23832335662d71f2993d4b019b5675d13c3d3040
 Author: jeremy <dpljer001@myuct.ac.za>
 Date: Tue Mar 27 15:58:26 2018 +0200

Test 5 working for AVL tree, all 4 graphs generating

commit 618bfc0d25b739e89eb4c646791f49b9b47b0ae2
 Author: jeremy <dpljer001@myuct.ac.za>
 Date: Tue Mar 27 15:18:13 2018 +0200

Test 5 working for BST

commit d34f52f40ea0bffd18eedd1b3b247c83d4a69995
 Author: jeremy <dpljer001@myuct.ac.za>
 Date: Mon Mar 26 22:41:57 2018 +0200

Opcount algorithm for insert comparisons working for BST, results writing to file for all tests

commit c5d72c4c8b41e295209c8ef14d88581e4fe58eec
 Author: jeremy <dpljer001@myuct.ac.za>
 Date: Mon Mar 26 19:26:25 2018 +0200

tests 2, 4, 5 and graph function all working, Makefile running everything correctly

commit c4b5872abcaa5536ad1d7b4edc9774a30811abc8
 Author: Jeremy du Plessis <dpljer001@myuct.ac.za>
 Date: Wed Mar 21 08:06:36 2018 +0000

Initial commit